

Week 7 - Regression Trees / x

https://canvas.northwestern.edu/courses/20673/assignments/98550

canvas COURSES & GROUPS GRADES CALENDAR NORTHWESTERN UNIVERSITY

2016SP_PRED_BUS_202-1_SEC57 2015 Spring

Home Syllabus Modules Discussions **Assignments** People Announcements Pages Files Quizzes Grades Course Reserves NU Canvas Resources

Week 7 - Regression Trees Assignment

Due Jun 3 by 1:29pm **Points** 15 **Submitting** a file upload

For the E-commerce data set, build a regression tree using rpart() in R. Use store_purchase_event_count as the dependent variable, and the other variables as the independent variables.

Generate the regression tree.

- Use the summary to view the diagnostics, and write a summary of the results
- Prune the tree to 10-20 significant nodes.
- Do you like or not like the tree? Why?
- Can you interpret the interactions in the tree?
- Print the trees, results of analysis, charts, and graphs, and your answers. Submit as Assignment 7.

Individual Assignment					Pts
Criteria	Ratings				
Content	Exceeds Expectations	Meets Expectations	Approaches Expectations	Does Not Meet Expectations	5 pts

Submission

✓ Turned In!
Jun 5 at 5:07pm (late)
[Submission Details](#) [Download Wk-7-Ec_R-Proj_Final_Upload.docx](#)

Comments:
 No Comments

[Re-submit Assignment](#)

Wk-7-Ec_R-Proj_.R

Rohit

Sat Jun 06 01:46:51 2015

Wk-7 EC - R- Project

```
ec <- read.csv("C:/STAT/Ec_R-Proj/ec.csv")
attach(ec)
library(caret)
```

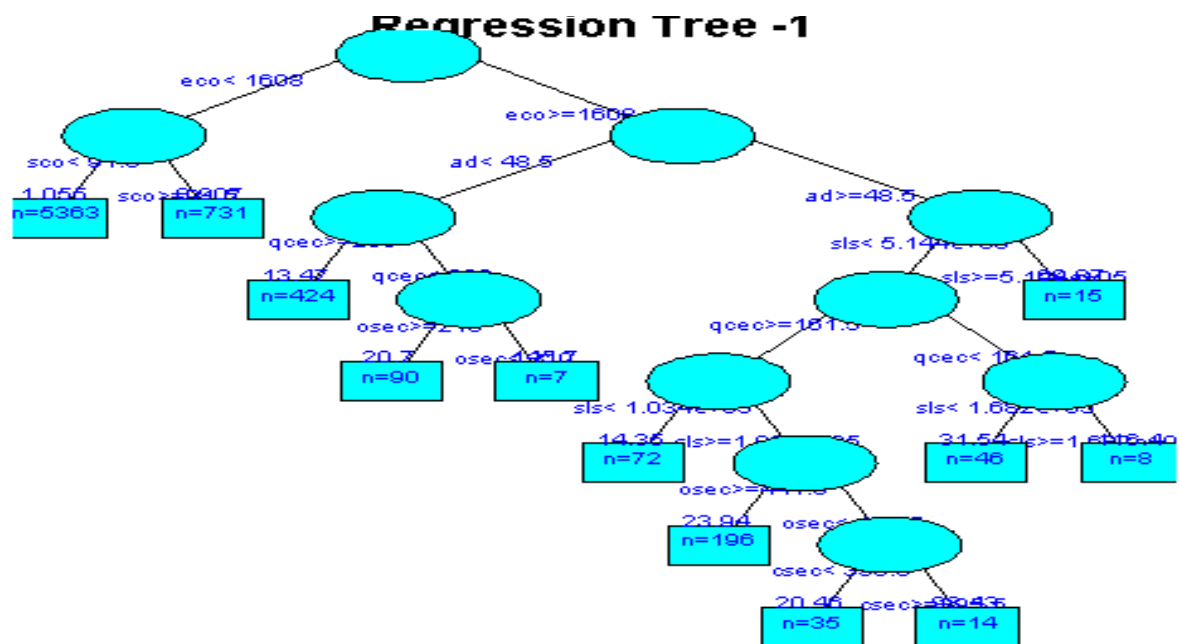
```
library(rpart)
library(rattle)
```

```
inTrain <- createDataPartition(y=ec$spec,p=0.7, list=FALSE)
trn <- ec[inTrain,]
tst <- ec[-inTrain,]
dim(trn); dim(tst)
```

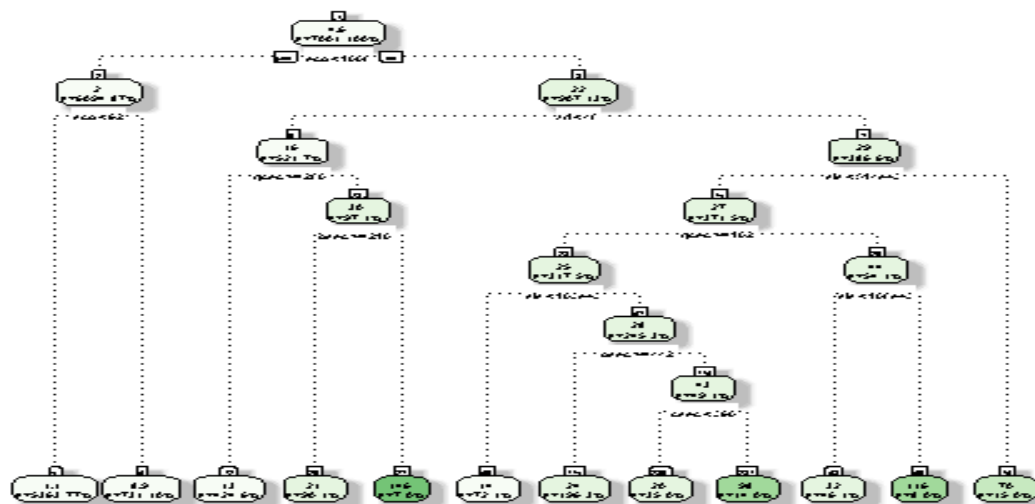
```
## [1] 7001 9
```

```
## [1] 2999 9
```

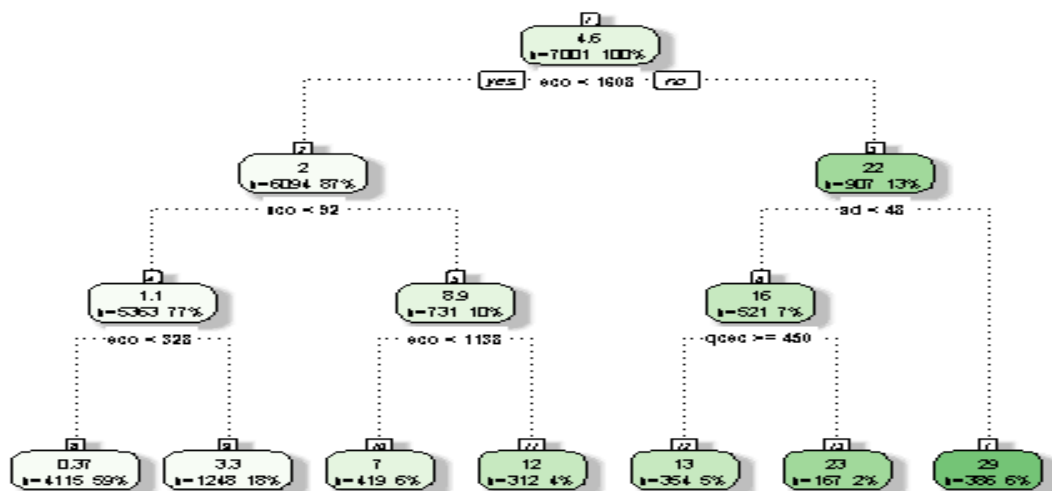
```
m<-rpart(spec~.,data=trn,method="anova")
par(mai=c(0.1,0.1,0.1,0.1))
plot(m,main="Regression Tree
-1",col=3,compress=TRUE,branch=0.2,uniform="TRUE")
text(m,cex=0.6,col=4,use.n=TRUE,fancy=TRUE,fwidth=0.4,fheight=0.4,bg=c(5)
)
```



fancyRpartPlot(m)



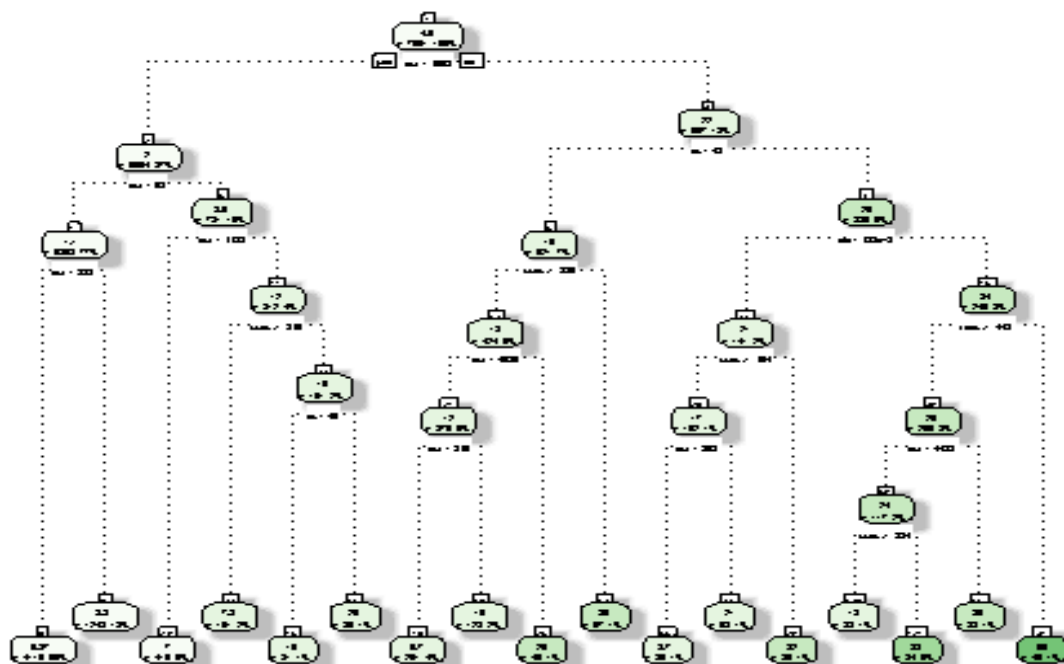
```
m_minsplit500cp0.001<-rpart(spec~.,data=trn,
method="anova",control=rpart.control(minsplit=500, cp=0.001))
library(rattle)
fancyRpartPlot(m_minsplit500cp0.001)
```



```

m_minsplit100cp0.001<-rpart(spec~.,data=trn,
method="anova",control=rpart.control(minsplit=100, cp=0.001))
fancyRpartPlot(m_minsplit100cp0.001)

```



```

TrnPred_m<-predict(m, newdata=trn)
head(TrnPred_m)

##      1      2      3      4      9     10
## 1.055379 1.055379 8.906977 20.457143 13.466981 1.055379

x<-as.matrix(TrnPred_m)
# As its a Regression problem - we need to Minimize Sum of Squared Errors
#
# Bagging - Bagging will have bias similar to the individual models
# but a reduced variance as we are averaging over individual models.

library(caret)
mBAG<-train(spec~.,method="treebag",data =trn)
print(mBAG$finalModel)

##
## Bagging regression trees with 25 bootstrap replications

## Random Forests

library(randomForest)

## randomForest 4.6-10
## Type rfNews() to see new features/changes/bug fixes.

rf <- randomForest(spec~., data=trn, mtry=6, importance = TRUE,ntree=10)
yhat_tst <- predict(rf, tst)
yhat_trn <- predict(rf, trn)
mean((yhat_tst - tst$spec)^2)

## [1] 308.2885

mean((yhat_trn - trn$spec)^2)

## [1] 58.06644

rf

##
## Call:
## randomForest(formula = spec ~ ., data = trn, mtry = 6, importance = TRUE,
ntree = 10)
##
##      Type of random forest: regression
##      Number of trees: 10
## No. of variables tried at each split: 6
##
##      Mean of squared residuals: 290.6614
##      % Var explained: -3.2


#Relative Importance
importance(rf)

##      IncNodePurity
## st      10660.97
## sls     285456.29
## sco      383617.35


```

```
## eco    253963.51
## csec   232026.07
## osec   294298.86
## qcec   191254.35
## ad     158977.83
```

```
varImpPlot(rf)
```



```
varImpPlot(rf, type=2, pch=11, col=4, cex=2, main="")
```



```
## Boosting ----
```

```
library(gbm)
```

```
## Loading required package: survival
```

```
##
```

```
## Attaching package: 'survival'
```

```
##
```

```
## The following object is masked from 'package:caret':
```

```
##
```

```
## cluster
```


```
##
```

```
## Loading required package: splines
```

```
## Loading required package: parallel
```

```
## Loaded gbm 2.1.1
```

```
boost <- gbm(spec~. , data=trn,distribution = 'gaussian', n.trees = 10,
interaction.depth = 4)
summary(boost)
```



```
##      var rel.inf
## eco  eco 50.436622
## sco  sco 17.220223
## csec csec 8.287838
## qcec qcec 8.249882
## osec osec 6.445820
## ad   ad 4.699561
## sls  sls 4.660054
## st   st 0.000000
```

as seen from Summary - eco eco 60.626854 and sco sco 10.632023 have high Relative Influence .

We plot them separately ...

```
par(mfrow=c(1,2))
plot(boost, i='eco')
plot(boost, i='sco')
```



```

boost.pred <- predict (boost, tst, n.trees=10)
mean((boost.pred - tst$spec)^2)

## [1] 276.3682

ctr <- trainControl(method = "cv", number = 10)
boost.caret <- train(spec~.,
trn,method='bstTree',preProc=c('center','scale'),trControl=ctr)

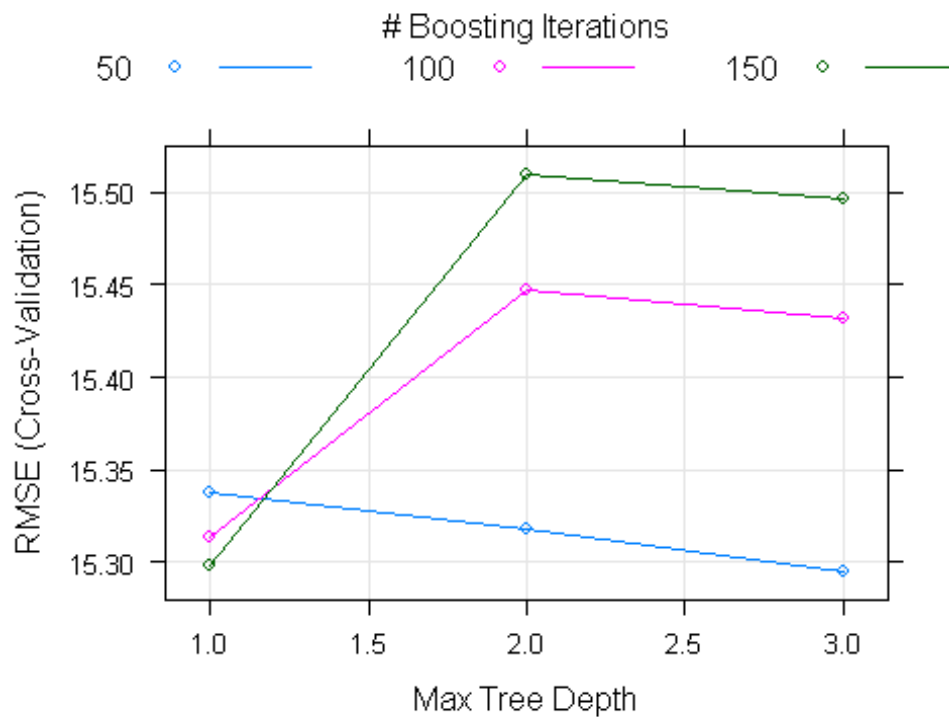
## Loading required package: bst

# yes
boost.caret

## Boosted Tree
##
## 7001 samples
## 8 predictor
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (10 fold)
##
## Summary of sample sizes: 6300, 6301, 6302, 6302, 6301, 6300, ...
##
## Resampling results across tuning parameters:
##
## maxdepth mstop RMSE Rsquared RMSE SD Rsquared SD
## 1 50 15.33740 0.2352452 4.248241 0.06221489
## 1 100 15.31355 0.2390788 4.226653 0.06484887
## 1 150 15.29802 0.2407234 4.229935 0.06566608
## 2 50 15.31806 0.2394373 4.169020 0.06417995
## 2 100 15.44771 0.2310972 4.093270 0.05860468
## 2 150 15.50997 0.2288551 4.050762 0.05757364
## 3 50 15.29485 0.2397998 4.179463 0.04864371
## 3 100 15.43241 0.2332206 4.045777 0.05172717
## 3 150 15.49668 0.2321351 4.017721 0.06180625
##
## Tuning parameter 'nu' was held constant at a value of 0.1
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were mstop = 50, maxdepth = 3 and nu
## = 0.1.

plot(boost.caret)

```



Here, with some better tuned parameters, we compare prediction - accuracy with random forests.??

```
boost.caret.pred <- predict(boost.caret, tst)
mean((boost.caret.pred - tst$spec)^2)
```

[1] 214.432

RF - > mean((boost.pred - tst\$spec)^2) == [1] 373.8108

boost.caret.pred - mean((boost.caret.pred - tst\$spec)^2) == [1] 311.0449

boost.caret.pred is better as - mean == 311.0449 is lower .